



OpenOKR

Your OKR coach, built in.

Open source, AI-native, self-hosted or in our cloud. The OKR method as executable rules, and two AI teammates that run the practice with you.

Most organisations do not fail at OKRs because their software was bad.

They fail because the practice never happened.

Somebody wrote objectives in a spreadsheet in January, nobody checked in by March, and the quarterly review was a meeting where everyone agreed things had gone reasonably well.

How OKRs are run today, and why the practice still dies

WHAT MOST START WITH

Spreadsheets

No cadence, no accountability, no quality bar.
Nobody is ever told anything. The file is opened twice a quarter.

WHAT MOST BUY

Conventional OKR trackers

A database with a progress bar. They store objectives faithfully and are entirely passive.
They will happily hold a badly written key result for a whole quarter without comment.

WHAT MOST FALL BACK ON

Consulting and training

The method is real, but it walks out of the door with the consultant. There is nothing to enforce it on a wet Tuesday in week six.

The practice that makes OKRs work has been well understood for twenty years. It just does not live in the software.

The failure pattern is always the same

- Objectives are written as project plans.
- Key results measure activity instead of impact.
- Nothing has a baseline, so nothing can be scored.
- Teams commit to more than they can deliver, and never record what they cut.
- Check-ins stop by week five.
- Blockers sit unowned.
- The quarterly review becomes a presentation rather than a diagnosis.
- The next quarter starts from a blank page, as if nothing had been learned.

Every one of those is **a rule that could have been enforced, a question that could have been asked at the right moment, and a person who could have been told.** That is the product.

01

What OpenOKR is

Two things make it different, and every design decision serves one of them.

The method is in the product, and the product is active

DIFFERENCE ONE

The method is in the product

The OKR practice canon is written as one authoritative specification and compiled into a pure library with no database, no network and no AI.

It holds the eight-phase cycle, the scoring and confidence bands, twenty quality rules with their word lists and coaching prompts, six publish gates, the alignment arithmetic, the KPI corridors, the blocker and root-cause taxonomies, both session agendas, and the closing diagnostic.

DIFFERENCE TWO

The product is active

Two agent members ship with every workspace. They are members, not features: they have names, they appear in feeds, they can be mentioned, and they are accountable.

They initiate, escalate and propose, in the browser, in Slack, Microsoft Teams, WhatsApp and Telegram, by email, and through whatever AI agent the user already runs.

One specification, compiled, running in four places at once

Because the rule engine is pure, the same code checks a draft in the browser, refuses a bad write on the server, drives the agents, and validates imported data.

In the browser

As the user types, on every
keystroke

On the server

Before any write is accepted

Inside the agents

So a nudge cites the same rule a
user sees

In the importer

So migrated data is judged by the
same bar

A conformance suite fails the build when the specification and the code disagree, so the product cannot drift away from the method. Every coaching message carries a rule key that resolves back to the specification, so a user can open the rule and argue with it.

Two AI teammates, with names, scopes and accountability



OKR Coach

Quality and practice

Reviews every draft against the twenty rules. Runs a nightly semantic sweep for duplicated metrics, better parents and hidden dependencies.

Speaks up when the not-doing list is empty, when nothing was cut at the capacity check, when a goal is reported on track but its key results have not moved in a month, and when the scores cluster suspiciously near perfect at the close.



OKR Champion

Rhythm and momentum

Reminds the champion before a check-in is due, on the day, and daily after. Escalates to the reviewer, the coordinator and then the sponsor, always visibly to the person being escalated past.

Runs the blocker clock, opens and closes the weekly session, assembles the digest, keeps the streak, watches every KPI corridor, and drafts the recovery objective when one drops out of range.

It all still works with AI switched off

The rules, nudges, escalations, gates, scores and the diagnostic are **deterministic code**, not prompts. AI adds drafting, rewriting, semantic judgement and natural language. **It never makes the decision.**

Sellable where AI is not

Regulated, air-gapped and AI-sceptical buyers get the whole practice. Continuous integration proves the product is whole with the provider off.

Propose by default

Every write an agent wants becomes a proposal in a human review queue. Direct writes need a narrow, explicit opt-in. Sandbox mode commits nothing.

Capped and least-privilege

Bindings on named spaces and goals only, never a workspace-wide grant. Every step metered, with a hard cost cap that halts a run mid-flight.

02

The product

What it actually looks like, screen by screen.

The Work Map is the front door

S-01

One company-wide tree. Every row carries the same contract: health, staleness, confidence, progress, champion and next step. An outdated badge overrides the last reported health, so nobody can hand-paint a goal green.

Northwind Group

FY2026 · Q3

Home

Review7

Inbox12

PRACTICE

CycleP5

Goals

KPIs2

Work

SPACES

Revenue

Product

Customer Success

People & Ops

YOUR AGENTS

C

OKR Coach

3 findings

Ch

OKR Champion

nudged 4

Work Map

Q Search goals, KPIs, people

36K

Ask AI

36J

+ New

AA

Phase 5 · Align and commit

2 gates unmet: capacity check, dependency confirmation

Publication deadline in 6 days

Northwind Group

Company-wide tree · 4 objectives · 14 key results · live

Company

My spaces

Revenue

Product

Q3 FY2026

Filter

GOAL / KEY RESULT	HEALTH	CONFIDENCE	PROGRESS	NEXT STEP	CHAMPION
OBJ Become the default supplier for mid-market retail in ASEAN	On track	High 0.8	64%	Confirm 2 dependencies	SL MR
KR Grow ASEAN mid-market ARR from RM 8.2m to RM 12.0m	On track	High 0.8	58%	RM 10.4m of RM 12.0m	JT
KR Raise win rate against Meridian from 22% to 35%	Caution	Med 0.6	39%	Blocker: pricing sign-off	PN
OBJ Make onboarding something customers finish without us	Outdated 11d	Med 0.5	41%	Escalated to sponsor	DK MR
KR Cut median time-to-first-value from 21 days to 7 days	Outdated	Low 0.3	29%	17 days · target 7	DK
INIT Self-serve setup wizard	Active	Med 0.5	7/10	Capacity: tight	RH
OBJ Turn support from a cost centre into a retention engine	Off track	Low 0.3	22%	Blocker 26h · escalated	FA SL
KPI Expansion revenue · drives net revenue retention	Unhealthy 43%	Recovering	43%	Recovery OKR live	FA
OBJ Ship the platform our 2027 roadmap depends on	On track	High 0.9	76%	Check-in due Friday	RH JT
KR Move platform p95 latency from 840ms to under 300ms	On track	High 0.9	81%	395ms · target 300ms	RH
KR Raise deploy frequency from 3 to 20 per week	On track	High 0.8	70%	15 of 20 per week	JT
KR Cut change-failure rate from 18% to under 5%	Caution	Med 0.6	52%	11% · target 5%	LW
INIT Migrate the ingest pipeline off the legacy queue	Active	High 0.8	17/20	Capacity: fits	LW

THE PRODUCT

12

A guided cycle that knows what is missing

Eight phases with computed completion, never self-reported ticks. Here drafting is locked because three of the seven input pack items are missing, and the block names them. A planning session without inputs produces objectives written from opinion.

Northwind Group
FY2026 - Q3

Home | Review | Inbox

PRACTICE

Cycle

Goals | KPIs | Work

SPACES

Revenue | Product | Customer Success | People & Ops

YOUR AGENTS

OKR Coach (idle) | **OKR Champion** (chasing 3)

Cycle / Q3 FY2026 planning

Search

Quarterly mode | Ask AI | AA

Phase 1 · Prepare | Drafting is blocked: the input pack is 4 of 7 complete | Publication deadline in 19 days

The cycle 2 of 8 done

- 0 · Annual strategy
The frame for FY2026
- 1 · Prepare**
Roles, dates, input pack
57%
- 2 · Diagnose
Scores, baselines, ranked issues
- 3 · Set direction
Priorities and the not-doing list
- 4 · Draft OKRs
Objectives and key results
- 5 · Align and commit
Cascade, capacity, six gates
- 6 · Run the cadence
Weekly, monthly, decisions
- 7 · Review and learn
Score, retro, feed forward

Phase 1 · Prepare Completion is computed, never self-reported

Drafting stays locked until the input pack is complete
A planning session without inputs produces objectives written from opinion. Missing: customer and market signal, capacity and constraints, the distribution confirmation.

Input pack 4 of 7

- ✓ Last cycle's scores and retrospective
Carried forward automatically from Q2 FY2026
- ✓ The annual strategy and annual OKRs
FY2026 frame · 3 strategies · 4 annual objectives
- ✓ KPI baselines and current health
18 KPIs · 2 outside their corridor
- ✓ Financial and operational plan
Uploaded by Sofia Lim · 2 days ago
- ✗ Customer and market signal
Assigned to Priya Nair · due in 3 days
- ✗ Capacity and known constraints
Assigned to Marc Reyes · due in 3 days
- ✗ Confirmed distribution to every participant
Distribute at least 3 working days before session one

Roles

- Sofia Lim**
Sponsor · decides and unblocks
- Marc Reyes**
Facilitator · guards the method

Session timeline All booked

Week	Activity	Due
Week 1	Diagnose and set direction	18 Aug
Week 2	Draft OKRs	25 Aug
Week 3	Align, capacity, commit	1 Sep

Facilitator guidance

KEY OUTPUT

A named sponsor, a named facilitator, every session in the calendar, and a complete pack in every participant's hands three working days before you start.

WATCH FOR

- A sponsor who delegates the decision. If they cannot attend, move the session.
- "We already know the numbers." Then it costs nothing to write them down.
- A pack sent the night before. People read it in the room and plan from opinion.

QUARTERLY MODE

The frame already exists. Phase 1 is a 30 minute confirmation, not a workshop. Reuse the annual sponsor unless the quarter changes owner.

OKR Champion Agent

I have asked Priya and Marc for the two missing items on Slack, and booked a reminder for tomorrow 09:00.

cycle.phase_blocked · CY-1

Twenty rules, running on every keystroke

S-09

The failing objective shows its verdicts inline. One key result measures activity, another has no baseline, and the whole set has no leading indicator. The strength meter and quality panel update live. Warnings never block typing.

Northwind Group

FY2026 · Q3

Home

Review

Inbox

PRACTICE

Cycle

Goals

KPIs

Work

SPACES

Revenue

Product

Customer Success

YOUR AGENTS

OKR Coach

reviewing

OKR Champion

idle

Cycle / Q3 FY2026 / Phase 4 · Draft OKRs

Search

Autosaved

Ask AI

JT

Phase 4 · Draft OKRs

3 checks failing across 2 objectives

Publication deadline in 11 days

Revenue · Q3 objectives

Every line is checked as you type. Warnings never block typing.

Draft from an ambition

+ Objective

Launch the new mobile ordering app in 3 markets

OBJ-1 · Outcome, not output

OBJ-2 · Metrics belong in the key results

AL-1 · Supports a bigger priority

OBJ-3 · Timebound

OBJ-4 · Owned

Hold 12 customer interviews in each market

increase

from 0

to 36

interviews

lagging

Rewrite

Improve customer satisfaction

increase

baseline missing

target missing

owner missing

Rewrite

Raise mobile share of orders from 12% to 30%

increase

from 12%

to 30%

30 Sep

lagging

Passing

KR-4 · Every key result here is lagging. You will only find out at the end of the cycle whether it worked. Add a leading indicator you can steer weekly.

Make mobile the way our customers prefer to reach us

100%

PN

SL

OBJ-1 ✓

OBJ-2 ✓

OBJ-3 ✓

OBJ-4 ✓

OBJ-5 ✓

AL-1 ✓ contributes to "Default supplier in ASEAN"

Increase NPS from 32 to 50

increase

from 32

to 50

NPS

lagging

Passing

Cut first-response time from 9h to 2h

reduce

from 9h

to 2h

hours

leading

Passing

Quality panel

3 must fix

LAUNCH THE NEW MOBILE ORDERING APP

OBJ-1 Outcome, not output

Your objective starts with a deliverable, not a destination.

AL-1 Supports a bigger priority

No parent and no stated contribution.

KR-5 Impact, not effort

"Hold 12 customer interviews" measures activity.

OBJ-2 Inspiring and directional

Contains digits. Metrics belong in the key results.

KR-4 Leading and lagging mix

All three key results are lagging.

KR-6 Ambitious but honest

Average confidence 0.92 at drafting reads as sandbagging.

OKR Coach

Agent

Objective 1 and objective 3 both move mobile order share. That is double counting: one team will get the credit twice.

Review finding

Dismiss

quality.conflict · semantic sweep · proposal, not a write

Coaching that is specific, and arguable

● Launch the new mobile ordering app in 3 markets |
2 rules firing

● OBJ-1 · Outcome, not output
FAIL

Your objective starts with a deliverable, not a destination. If we do it and nothing changes, did we succeed? Rewrite around the change you want.

Why it matters. An objective you can complete without anything improving is a task list. The launch is a means, not the end.

WEAK
Launch the new mobile app by end of Q3

STRONG
Make mobile the way our customers prefer to reach us

✨ Rewrite this for me
See the rule in METHOD
Dismiss

EVERY VERDICT OPENS INTO THE RULE

The coaching prompt, why it matters, a weak-versus-strong pair, and a rewrite the user applies or dismisses. Never a bare error string, and never a mysterious red dot.

SOME OF WHAT IT CATCHES

- An objective starting with "launch". If we launch it and nothing changes, did we succeed?
- A key result with a target but no baseline. Without the "from", movement cannot be proved.
- "Hold twelve customer interviews." That is activity. What are they for? Measure that.
- Every key result lagging. You only find out at the end. Add a leading indicator.
- Average confidence above ninety percent at drafting. That is sandbagging, not a stretch.

Six publish gates, enforced server-side

- 1 Every objective has a title, a named champion and a named reviewer.
- 2 Every key result passes its checks.
- 3 Alignment is mapped. Each objective states what it contributes to.
- 4 Every dependency is confirmed, or logged with a named risk owner.
- 5 Capacity is checked, nothing is left exceeding, and the cuts are recorded.
- 6 A publication date is set before day one of the cycle.

Gate five is the one nobody has had

Capacity is read from the initiatives actually planned against each key result, not self-declared.

If one exceeds and no cut is recorded, the cycle cannot be published.

A plan where nothing was cut is a plan that has not been made.

The publish control states its reason, never sits silently inert

Two gates are red. Capacity exceeds on one key result with nothing recorded as cut, and a dependency on Finance is unconfirmed. The Coach names the consequence rather than just the rule.

Northwind Group

FY2026 · Q3

Home

Review

Inbox

PRACTICE

Cycle

Goals

KPIs

Work

SPACES

Revenue

Product

Customer Success

YOUR AGENTS

OKR Coach

OKR Champion

Cycle / Q3 FY2026 / Phase 5 · Align and commit

Search

Ask AI

MR

Phase 5 · Align and commit

2 of 6 gates unmet

Publication deadline in 6 days

Capacity check

Read from the initiatives actually planned against each key result

KEY RESULT	MAIN INITIATIVES	TEAM LOAD	VERDICT
Grow ASEAN mid-market ARR to RM 12.0m	Partner programme, outbound pod, pricing refresh	3 of 4 FTE	Fits
Raise win rate against Meridian to 35%	Battlecards, competitive pricing sign-off	2 of 2 FTE	Tight
Cut median time-to-first-value to 7 days	Self-serve wizard, template library, in-app guide, migration tooling	6 of 4 FTE	Exceeds
Increase NPS from 32 to 50	Response-time programme	1 of 2 FTE	Fits

What did you cut?

One key result exceeds capacity and this field is empty. A plan where nothing was cut is a plan that has not been made. Required before gate 5 turns green.

Name the initiative you are dropping, and who was told...

Dependency register

1 unconfirmed

KEY RESULT	NEEDS	FROM	STATE	RISK OWNER
Cut time-to-first-value to 7 days	Migration API	Platform	Confirmed	—
Raise win rate to 35%	Pricing sign-off	Finance	Unconfirmed	MR Marc Reyes
Increase NPS from 32 to 50	Support headcount	People & Ops	Confirmed	—

Publish gates

4 of 6

1 · Every objective owned and reviewed

7 objectives · champion and reviewer set on all

2 · Every key result passes its checks

23 key results · all seven checks green

3 · Alignment is mapped

Every objective states what it contributes to

4 · Every dependency confirmed or risk-owned

Pricing sign-off from Finance is unconfirmed. [Open the register](#)

5 · Capacity checked, nothing exceeding

1 key result exceeds and no cut is recorded. [Fix the capacity check](#)

6 · Publication date before day one

Publishes 29 Sep · cycle starts 1 Oct

Publish the cycle

Disabled until gates 4 and 5 are green. The reason is always stated, never silently inert.

OKR Coach


Agent

Nothing has been cut anywhere in this cycle, yet one key result is over capacity. In my experience that means the cut happens in week six instead, without a decision.

quality.no_cuts · CY-6 · sent to Marc Reyes on Slack


Alignment is contribution, not copying

Solid connectors are contribution, dashed are dependencies both teams know about. The orphan goal is flagged, the health score names every gap with its penalty, and the Coach adds what structure alone cannot see.



When a KPI leaves its corridor, the product drafts the fix

Expansion revenue sits at forty-three percent of target. It was forty percent when the recovery launched and the recovery is forty-nine percent complete, so effective health reads sixty-five: the fix shows before the lagging number catches up. One key result per leading child driver.



Fifteen minutes, four steps, run by the product

S-22

Team votes reveal together so nobody anchors on the champion. Every low score becomes a typed blocker with a named owner and one action inside twenty-four hours. Last week's commitments close honestly, including the one that did not land.

Northwind Group

Revenue space

Home

Review

Inbox

PRACTICE

Cycle

Goals

KPIs

Work

SPACES

Revenue

Product

Customer Success

RUNNING THIS SESSION

Ch OKR Champion

live

Revenue / Weekly check-in · Friday 14 Aug

Search

96K

Session live · 08:41 elapsed

JT PN LW SL Ch

Weekly check-in

Fifteen minutes, four steps. It produces decisions, not status.

1 Confidence round

4 of 6 scored

2 Diagnose the lows

Type, owner, 24h action

3 Commitments

Close last, set this

4 Digest

Generated, then posted

12 WEEK CONFIDENCE

Falling four weeks running.

Key results

Score 0.0 to 1.0

0.8 Grow ASEAN mid-market ARR to RM 12.0m

JT

0.3 Raise win rate against Meridian from 22% to 35%

PN

0.7 Increase partner-sourced pipeline to RM 5.0m

LW

0.35 Cut sales cycle from 74 to 45 days

JT

– Raise average deal size from RM 42k to RM 60k

PN

– Lift proposal win rate on renewals to 85%

LW

NOW SCORING

Raise win rate against Meridian from 22% to 35%

Champion Priya Nair · now 26% · leading

9 week streak

Skip step

Continue to blockers

Rhythm streak

9 weeks

A skipped week breaks it. Nothing else does.

Open blockers

1 overdue

26h Pricing sign-off from Finance

Dependency · Marc Reyes · escalated to sponsor

20h Migration API scope unclear

Clarity · Jun Tan

6h Two analysts pulled to the audit

Resource · Lena Wong

Last week's commitments

No negotiation

Ship the Meridian battlecard to all AEs

Priya Nair · delivered

Book the top 20 renewal calls

Lena Wong · delivered

Agree the discount ladder with Finance

Jun Tan · not delivered, second week

Step 2 · Diagnose what is low

Every score below 0.4 needs a type, an owner and one action within 24 hours

0.3 Raise win rate against Meridian

At 0.3, escalated today

TYPE Resource Dependency Clarity Priority conflict External

Dependency: another team owes you something before you can move. Name them and the date.

OWNER MR Marc Reyes

Get pricing sign-off from Finance in Monday's exec slot

20h left

0.35 Cut sales cycle from 74 to 45 days

Clarity JT

Action set

Sixty minutes, eleven timed stages, one honest diagnosis

S-24

Scores stay hidden until the room reveals them together. Every key result below 0.7 gets exactly one cause. Five process-health statements are scored anonymously, and the lowest becomes next quarter's process priority.

Northwind Group

Q2 FY2026 close

Home

Review

Inbox

PRACTICE

CycleP7

Goals

KPIs

Work

SPACES

Revenue

Product

Customer Success

PREPARED THIS SESSION

OKR Champion

OKR Coach

pack ready

watching

Cycle / Q2 FY2026 / Quarterly review

Stage 7 of 11 · Root cause and diagnostic

4:12 of 5:00

+ 1 min

SL

JD

DI

FA

PH

MR

Export minutes

OPEN

Open and check-in5

REVIEW

Score the key results12

Objective narratives9

Recognition and wins3

RETRO

Team retro7

Management retro3

Root cause and diagnostic5

OKR process health3

RESET

Keep, modify or abandon5

Learnings and next drafts4

Decisions and actions4

The diagnostic

Cycle score 0.58 · rhythm 3.8 of 5

A strategy or OKR-quality problem, not a people problem.

You scored below 0.7 while running the rhythm well: nine weeks of check-ins, blockers named and cleared, commitments closed honestly. The team did the practice and still missed. That points at the key results themselves, or the strategy behind them. Fix the key results before you push the team.

Computed from your data. It is the same answer with AI switched off.

See how this is computed

One honest cause for every key result under 0.7

3 of 5 assigned

Raise win rate against Meridian from 22% to 35%

0.31

Ambition too high

Wrong key result

Blocked by a dependency

Capacity

Priority shifted

External change

Too many OKRs

No owner or cadence

Win rate was the wrong measure. We only compete with Meridian in one segment, so the number moved on deal mix, not on anything we did.

Cut median time-to-first-value from 21 to 7 days

0.42

Ambition too high

Wrong key result

Blocked by a dependency

Capacity

Priority shifted

External change

Too many OKRs

No owner or cadence

Increase partner-sourced pipeline to RM 5.0m

Cause needed0.55

Look for the system, not the person. Ask why until it stops being a symptom. METHOD \$8.4

Scorecard

Cycle 0.58

Default supplier in ASEAN

0.64

Grow mid-market ARR

0.82

Raise win rate vs Meridian

0.31

Partner-sourced pipeline

0.55

Onboarding without us

0.47

Time-to-first-value

0.42

Onboarding completion

0.52

Each objective score stayed hidden until the room revealed it together. Nobody anchored.

Process health

Anonymous - 6 of 6

OKRs stayed visible and drove decisions

4.2

A real check-in cadence, not a status report

4.5

Key results measured outcomes, not activity

2.6

Few enough OKRs that focus was possible

3.8

We said so early when things went off track

3.9

Lowest statement becomes next cycle's process priority

"Key results measured outcomes, not activity" scored 2.6. It is already placed in Phase 3 of Q3 planning.

A coach that only lives in a browser tab is not active

Every channel is two-way: nudges out, and real work in. A fully conversational WhatsApp check-in captures a typed blocker and its owner. Every message names the rule that sent it and where it sits on the escalation ladder.

The coach reaches you where you already work

Every message names the rule that sent it, offers a snooze and a way to change channel, and is recorded as a nudge with its escalation step. One per subject per person per day, unless it escalates.

Slack · #revenue

Ch OKR Champion APP 09:00

Morning Dinesh. Your check-in on **Make onboarding something customers finish without us** was due Friday. It is the third reminder, so I have let Marc know.

Cut median time-to-first-value from 21 to 7 days

Now	17 days
Last confidence	0.3 low
Last check-in	11 days ago
Escalation	Step 3 of 5 · reviewer

Check in

Log a blocker

Snooze 1 day

checkin.overdue · why am I getting this?

snooze · change my channel

Ch OKR Champion APP 16:45

Revenue weekly digest · week 6

On track	3 of 6 key results
At risk	Win rate vs Meridian · 0.4
Blockers	1 open, 26h, escalated
Commitments	2 delivered, 1 not
Streak	9 weeks

Open the digest

Post to #general

digest.weekly · after the session closes

Email nudges, digests, the daily summary, one-click check-in links

Microsoft Teams · Revenue

C OKR Coach BOT 14:20

Two objectives in this cycle both move mobile order share. One quarter, two teams, the same number.

DUPLICATE METRIC · AL-2

Revenue and Partnerships both count partner-sourced ARR

Objective A Win the accounts...

Objective B Make partners our...

Overlap RM 2.1m of target

Proposed Partnerships owns it

Review the proposal

Dismiss

A proposal, not a write. Nothing changes until a human applies it.

Ch OKR Champion BOT 15:30

The weekly session is open. Six key results to score, three blockers to clear.

Open the session

I cannot make it

session.open · at the scheduled start · quiet hours respected in your timezone

Browser everything

WhatsApp · Business

Hi Priya. Weekly check-in for **Raise win rate against Meridian**. Reply with the current number and your confidence 0 to 1.

09:00

26% and 0.4

09:04 ✓✓

Logged. 0.4 is below the line, so this needs a blocker type. Reply with a number:

1 Resource · 2 Dependency

3 Clarity · 4 Priority conflict · 5 External

09:04

2

09:05 ✓✓

Dependency noted. Who owns clearing it, and what is the one action in the next 24 hours?

09:05

Marc. Pricing sign-off in Monday's exec slot

09:06 ✓✓

Done. Marc has the blocker with a 24 hour clock, and your check-in is published. **Nothing else is due from you today.**

09:06

Your own AI agent Claude, ChatGPT, Cursor or a custom agent, acting as you, within your permissions

Telegram · OpenOKR bot

Blocker escalation

"Pricing sign-off from Finance" has been open **26 hours** with no action logged. The clock is 24.

Key result: Raise win rate vs Meridian · Owner: Marc Reyes · Confidence 0.4

I have cleared it

Reassign

Escalate to Sofia (sponsor)

Snooze 4h

See the rule

17:00

Escalate to Sofia

17:02

Sofia has been told, and Marc can see that I escalated.

blocker.overdue · coordinator at 24h, sponsor at 48h

17:02

Daily summary · Friday digest.daily

2 check-ins due · 1 acknowledgement owed · 1 blocker you own · weekly session at 16:00

Open my review inbox

08:00

No chat provider a workspace with none configured still works completely

THE PRODUCT

24

Notifications say what happened. This says what you owe.

S-02

Server-computed, overdue first, driving a live badge. Agent proposals queue here too. A snooze quietens the message and never hides the obligation.

Northwind Group

FY2026 · Q3

Home

Review7

Inbox12

PRACTICE

Cycle

Goals

KPIs

Work

SPACES

Revenue

Product

Customer Success

People & Ops

YOUR AGENTS

OKR Coach

2 proposals

OKR Champion

nudged you 2x

Review

Q Search

96K

Ask AI

96J

+ New

SL

What you owe

Server-computed, overdue first. Notifications say what happened. This says what you owe.

OVERDUE 3

Post your check-in on "Make onboarding something customers finish without us"

Champion · weekly · Product space · confidence was 0.3 last week

Overdue by 4 days

Check in

Clear the blocker "Pricing sign-off from Finance"

Dependency · Raise win rate vs Meridian · escalated to you as sponsor

26h on a 24h clock

Take action

Close last week's commitment: "Agree the discount ladder with Finance"

Commitment · Revenue space · second week open

Overdue by 8 days

Delivered / not

DUE TODAY 2

Acknowledge 3 check-ins from your team

Reviewer · Jun Tan, Lena Wong, Priya Nair · published yesterday

Today

Review all 3

Run the weekly session for Revenue

Coordinator · 6 key results to score · 3 open blockers · streak 9

16:00 today

Start session

AGENT PROPOSALS WAITING ON YOU 2

Coach proposes: split the duplicated partner-ARR metric between two objectives

quality.conflict · semantic sweep · would edit 2 key results · nothing changes until you apply

Dismiss

Review

Champion drafted a recovery objective for "Expansion revenue"

kpi.recovery_proposed · unhealthy for two periods · 2 key results, one per leading child driver

Dismiss

Review

THIS WEEK 4

Post your check-in on "Become the default supplier for mid-market retail in ASEAN"

Champion · weekly · Company · on track, 64%

Friday

Check in

Acknowledge Dinesh Kumar's check-in on onboarding

Reviewer · Product space · confidence dropped 0.5 to 0.3

Thursday

Review

2 more this week

OVERDUE 3

TODAY 2

THIS WEEK 4

Why you got nudged

Provenance

Ch "Your check-in on Make onboarding... was due Friday. It is the third reminder, so I have let Marc know."

checkin.overdue · sent 09:00 on Slack · escalation step 3 of 5

See the rule

Snooze

Change my channel

A snooze quietens the message. It never hides what you owe here.

Escalation ladder

1 Due date reached · the champion

done

2 1 day overdue · the champion again

done

3 3 days overdue or grace · the reviewer

now

4 7 days overdue · the coordinator

5 14 days overdue · the sponsor

Every step is visible to the person being escalated past. Nothing happens behind anyone's back.

Your week

Check-ins on time

7 of 9

Acknowledgements closed

12 of 12

Median blocker to action

9h

Target is under 24 hours.

THE PRODUCT

25

03

Why it wins

What no other OKR tool does, and what each of you gets.

Five things no other OKR tool does

1

The method is executable, not documentation

Competitors ship templates and help articles. We ship rules that run on every keystroke and refuse a bad publish. A customer cannot quietly stop doing the practice.

2

The software initiates

Every other tool waits to be visited. Ours arrives in Slack on Friday morning, escalates the blocker that aged past its clock, and tells the sponsor early.

3

It works with AI switched off

The deterministic core makes it sellable into regulated, air-gapped and AI-sceptical buyers who would reject an LLM-dependent product outright.

4

It diagnoses, not just reports

At the close it says whether a missed quarter was a cadence problem or a strategy problem. Every executive asks that. No tracker answers it.

5

Agent-native to the core

Everything a human can do, an agent can do, through one permission-checked contract. The customer's own Claude or ChatGPT is a first-class user, not a scraping workaround.

What each of you gets

A methodology institute

The practice becomes enforceable after the engagement ends.

Every rule, band, corridor, taxonomy and agenda lives in one versioned, auditable specification, separable from the code. Coaching prompts and rule names are translatable strings.

You can see exactly what the product teaches, and disagree with it in a specific place rather than in general.

An investor

A category with real budget, sold today mostly as passive databases.

Open source removes funnel cost and makes the enterprise security review a non-event. AGPL blocks a hyperscaler reselling it. The cloud sells operation, so there is one product, one release and no feature-gate resentment.

The agent layer is the moat: accumulated method, not a copyable prompt.

An early customer

Setup to first check-in in under fifteen minutes.

Data on your own servers, in your own country, with no seat limit and no feature gate. Existing goals import from a spreadsheet with an AI-assisted mapper and a dry run.

The whole workspace exports whenever you want. Nothing creates lock-in, which is exactly why it is easy to say yes to.

The outcomes we are designing for

Leading indicators, instrumented in the product from day one.

< 15 min

From install to first published check-in

70%+

Of due check-ins submitted on time

> 75%

Median OKR strength score at publication

< 24 h

Median blocker logged to next action taken

6+ weeks

Rhythm streak in half or more of active spaces

Lagging: active instances after six months, organisations completing a full cycle end to end, and organisations running an agent unattended every week.

04

The business

Scope, licence, status, and what we are asking for.

Everything in version one, across six pillars

Nothing is gated behind a paid tier. Self-host gets every capability, with no seat limit.

A The OKR core

Cycles and the guided workflow · goals and direction-aware key results · alignment with dependencies and the health score · KPIs, driver trees and formulas · health corridors and recovery OKRs · check-ins with snapshots and team voting · scorecard

B The rhythm

Weekly check-in session · the five-type blocker taxonomy on a 24-hour clock · commitments · monthly review and decision log · the eleven-stage quarterly review · mid-cycle calibration · cadence and staleness · the review inbox and digests

C The work

Initiatives with a capacity verdict · tasks and the key-result-linked board · rich documents with version history · file attachments. Deliberately OKR-shaped, not a project management suite

D Coaching and AI

The Draft Coach rule engine · the OKR Coach and OKR Champion · agent governance with least privilege, proposals and hard caps · the nudge engine · assists everywhere · the copilot · bring your own provider, including local models

E Channels and reach

Browser · email · Slack · Microsoft Teams · WhatsApp Business · Telegram · any external AI agent acting as its user. Every inbound message runs the same permission checks as a click

F Platform

Spaces · people and org chart · relationship-based access · comments, mentions and notifications · activity feed · search and command palette · admin and audit · encrypted export and import · spreadsheet and FlowyTeam importers · accessibility

Fully specified, and sequenced so the coach is not last

Eight phases, 104 scoped tasks, each with acceptance criteria and a test plan.

1 Foundation

Monorepo, database with the tenant floor, adapter ports and the outbox, authentication, the single write pipeline, Compose and Helm

3 The OKR core

Cycles and the guided workflow, goals and key results, scoring, cadence, check-ins, alignment, KPIs with recovery, the Work Map

5 Reach

Slack, Teams, WhatsApp, Telegram, the external agent server, initiatives, tasks and the board, documents, search

7 Hardening

Performance at scale, load and soak testing, the security review, the accessibility audit, observability

2 Platform and agent spine

Access model, people, notifications, design system, and the AI foundation: provider port, keys, metering, the agent runtime

4 The coaching layer

The method library, the Draft Coach, both agents, the trigger catalogue, the weekly and quarterly sessions, the copilot

6 Data

The spreadsheet importer with the AI mapper, the FlowyTeam importer, workspace export and import, backups with restore drills

8 Cloud and launch

Tenant provisioning, the operator console, single sign-on and directory sync, the air-gap guide, documentation, launch

The AI and agent foundation lands in phase two so the coaching layer ships in phase four, beside the OKR core. **An OKR tool where the coach arrives last is just another tracker.**

The risk we are managing, openly

A coaching engine that produces false positives is worse than no coaching at all, because people learn to dismiss it.

What we are doing about it

The warn-versus-fail line on all twenty rules will be tuned against a corpus of real, anonymised OKRs before launch. Assembling that corpus is a named deliverable of the coaching phase design gate, not an afterthought.

Where a partner adds the most value

This is the single place where methodology expertise changes the product most. A partner who can challenge a threshold, or contribute real OKRs to tune against, moves the quality of every coaching message in the product.

Three different conversations

A methodology institute

WHAT WE ASK

Review the method specification, and challenge any rule, threshold, band or agenda you think is wrong. Help us assemble the tuning corpus.

WHAT YOU GET

Your practice becomes enforceable software that keeps working after the engagement ends, with every rule attributable, versioned and arguable.

An investor

WHAT WE ASK

A conversation about the category, the open source route to market, and the plan to first revenue.

WHAT YOU GET

A fully specified product with an unusually defensible core, in a market currently served by passive databases.

An early customer

WHAT WE ASK

Run a real quarter on it, with us, and tell us where the coaching is wrong.

WHAT YOU GET

Free use for the pilot, direct influence on the rules and the roadmap, and no lock-in of any kind because the whole workspace exports at any time.

The practice is the product.

Every rule, threshold, band, corridor, taxonomy and agenda in one specification, compiled into code, enforced at the point of writing, and worked by two teammates who initiate.

The full planning set behind this deck is available on request: requirements, architecture, the method canon, the schema, forty screen specifications and 104 tasks.